Dario Altman

AI Infrastructure Engineer

✉ dario.altman@example.com | example.com/dario-altman
✆ +1 (202) 555-0148 | example.com/dario-code

PROFESSIONAL SUMMARY

AI infrastructure engineer building reliable, efficient model-serving systems from prototype through production. Experienced in inference optimization, distributed platforms, and practical evaluation, with a focus on making machine learning services observable and easy for product teams to operate.

EDUCATION

B.S. in Computer Engineering | Lakehurst University — Sep 2014 – May 2018 | GPA: 3.8/4.0

SKILLS

Languages	Python, Go, SQL, Bash, TypeScript
AI & Inference	PyTorch, ONNX Runtime, vLLM, model evaluation, quantization
Infrastructure	Kubernetes, Docker, Terraform, gRPC, Linux
Observability	Prometheus, Grafana, OpenTelemetry, load testing, incident response

PROJECTS

Inference Bench | Reproducible model-serving benchmark suite | example.com/bench

- Built a configurable harness to compare throughput, tail latency, and memory use across inference backends, with realistic request shapes and repeatable traffic patterns under load.
- Added repeatable workload fixtures and reports that make performance changes easy to review before deployment and help teams compare results across different hardware configurations.
- Documented setup, interpretation, and trade-offs so engineers can reproduce benchmarks on a single GPU or a shared cluster and select a suitable backend for each workload in production.

Signal Canvas | Model-serving observability dashboard | example.com/signal-canvas

- Created a dashboard linking request traces, queue depth, GPU utilization, and error rates for deployed models, making it easier to trace slow responses to resource constraints during incidents.
- Implemented alert examples and a lightweight replay tool to help teams investigate latency regressions, reproduce failed requests, and compare behavior across releases and production environments.
- Published a sample deployment with synthetic traffic and practical runbooks for local exploration, onboarding, and common troubleshooting exercises across serving environments with teammates.

EXPERIENCE

AI Platform Engineer | Lattice Harbor | Remote — Jun 2021 Present

- Designed a shared inference platform for language and embedding models, giving product teams a consistent API, deployment templates, and a safer path from prototype to production.
- Introduced dynamic batching and autoscaling based on queue pressure, reducing median serving cost by 31% across representative workloads while maintaining latency targets for customers.
- Established latency budgets, tracing, and release checks that cut high-severity serving incidents by 40% over twelve months and made investigations faster for on-call engineers across teams.

Software Engineer | Emberfield Cloud | Seattle, WA — Jul 2018 – May 2021

- Developed containerized services and deployment tooling for internal data products, improving release consistency across four engineering teams and simplifying operational handoffs for each release.
- Built a load-testing pipeline and capacity reports that revealed bottlenecks before launches, reduced manual performance checks, and gave teams a shared baseline for planning and scaling.
- Partnered with data scientists to package prediction models as monitored services, with documented ownership, deployment procedures, and tested rollback paths for each release in production.

EXTRA-CURRICULAR ACTIVITIES & ACHIEVEMENTS

- Maintain an open-source model evaluation toolkit and mentor engineers through a local machine learning reading group, organizing practical sessions on serving and reliability for practitioners.